

Data Augmentation To Address Overfitting In Flower Classification CNN

In this notebook we will build a CNN to classify flower images. We will also see how our model overfits and how overfitting can be addressed using data augmentation. Data augmentation is a process of generating new training samples from current training dataset using transformations such as zoom, rotations, change in contrast etc

Credits: I used tensorflow offical tutorial:

<https://www.tensorflow.org/tutorials/images/classification> as a reference and made bunch of changes to make it simpler

In below image, 4 new training samples are generated from original sample using different transformations

```
In [2]: import matplotlib.pyplot as plt
import numpy as np
import cv2
import os
import PIL
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
```

We will download flowers dataset from google website and store it locally. In below call it downloads the zip file (.tgz) in cache_dir which is . meaning the current folder

Load flowers dataset

```
In [3]: dataset_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz"
data_dir = tf.keras.utils.get_file('flower_photos', origin=dataset_url, cache_dir='.')
# cache_dir indicates where to download data. I specified . which means current directory
# untar true will unzip it
```

```
In [4]: data_dir
```

```
Out[4]: '.\datasets\flower_photos'
```

```
In [5]: import pathlib
data_dir = pathlib.Path(data_dir)
data_dir
```

```
Out[5]: WindowsPath('datasets/flower_photos')
```

```
In [6]: list(data_dir.glob('*/*.jpg'))[:5]
```

```
Out[6]: [WindowsPath('datasets/flower_photos/daisy/100080576_f52e8ee070_n.jpg'),
          WindowsPath('datasets/flower_photos/daisy/10140303196_b88d3d6cec.jpg'),
```

```
WindowsPath('datasets/flower_photos/daisy/10172379554_b296050f82_n.jpg'),
WindowsPath('datasets/flower_photos/daisy/10172567486_2748826a8b.jpg'),
WindowsPath('datasets/flower_photos/daisy/10172636503_21bededa75_n.jpg')]
```

In [7]:

```
image_count = len(list(data_dir.glob('/*/*.jpg')))
print(image_count)
```

3670

In [8]:

```
roses = list(data_dir.glob('roses/*'))
roses[:5]
```

Out[8]:

```
[WindowsPath('datasets/flower_photos/roses/10090824183_d02c613f10_m.jpg'),
 WindowsPath('datasets/flower_photos/roses/102501987_3cdb8e5394_n.jpg'),
 WindowsPath('datasets/flower_photos/roses/10503217854_e66a804309.jpg'),
 WindowsPath('datasets/flower_photos/roses/10894627425_ec76bbc757_n.jpg'),
 WindowsPath('datasets/flower_photos/roses/110472418_87b6a3aa98_m.jpg')]
```

In [9]:

```
PIL.Image.open(str(roses[1]))
```



In [10]:

```
tulips = list(data_dir.glob('tulips/*'))
PIL.Image.open(str(tulips[0]))
```



Read flowers images from disk into numpy array using opencv

In [9]:

```
flowers_images_dict = {
    'roses': list(data_dir.glob('roses/*')),
    'daisy': list(data_dir.glob('daisy/*')),
```

```
'dandelion': list(data_dir.glob('dandelion/*')),
'sunflowers': list(data_dir.glob('sunflowers/*')),
'tulips': list(data_dir.glob('tulips/*')),
}
```

In [10]:

```
flowers_labels_dict = {
    'roses': 0,
    'daisy': 1,
    'dandelion': 2,
    'sunflowers': 3,
    'tulips': 4,
}
```

In [13]:

```
flowers_images_dict['roses'][:5]
```

Out[13]:

```
[WindowsPath('datasets/flower_photos/roses/10090824183_d02c613f10_m.jpg'),
 WindowsPath('datasets/flower_photos/roses/102501987_3cdb8e5394_n.jpg'),
 WindowsPath('datasets/flower_photos/roses/10503217854_e66a804309.jpg'),
 WindowsPath('datasets/flower_photos/roses/10894627425_ec76bbc757_n.jpg'),
 WindowsPath('datasets/flower_photos/roses/110472418_87b6a3aa98_m.jpg')]
```

In [14]:

```
str(flowers_images_dict['roses'][0])
```

Out[14]:

```
'datasets\\flower_photos\\roses\\10090824183_d02c613f10_m.jpg'
```

In [11]:

```
img = cv2.imread(str(flowers_images_dict['roses'][0])) ## str because of cv2 transfo
```

In [12]:

```
img.shape
```

Out[12]:

```
(240, 179, 3)
```

In [17]:

```
cv2.resize(img,(180,180)).shape # o have the same size
```

Out[17]:

```
(180, 180, 3)
```

In [13]:

```
X, y = [], []

for flower_name, images in flowers_images_dict.items():
    for image in images:
        img = cv2.imread(str(image))
        resized_img = cv2.resize(img,(180,180))
        X.append(resized_img)
        y.append(flowers_labels_dict[flower_name])
```

In [14]:

```
X = np.array(X)
y = np.array(y)
```

Train test split

In []:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

Preprocessing: scale images

```
In [ ]: X_train_scaled = X_train / 255  
X_test_scaled = X_test / 255
```

Build convolutional neural network and train it

```
In [22]: num_classes = 5  
model = Sequential([  
    layers.Conv2D(16, 3, padding='same', activation='relu'),  
    layers.MaxPooling2D(),  
    layers.Conv2D(32, 3, padding='same', activation='relu'),  
    layers.MaxPooling2D(),  
    layers.Conv2D(64, 3, padding='same', activation='relu'),  
    layers.MaxPooling2D(),  
    layers.Flatten(),  
    layers.Dense(128, activation='relu'),  
    layers.Dense(num_classes)  
])  
  
model.compile(optimizer='adam',  
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              metrics=['accuracy'])  
  
model.fit(X_train_scaled, y_train, epochs=30)
```

```
Epoch 1/30  
86/86 [=====] - 162s 1s/step - loss: 1.3306 - accuracy: 0.4419  
Epoch 2/30  
86/86 [=====] - 94s 1s/step - loss: 1.0132 - accuracy: 0.6119  
Epoch 3/30  
86/86 [=====] - 94s 1s/step - loss: 0.8252 - accuracy: 0.6893  
Epoch 4/30  
86/86 [=====] - 94s 1s/step - loss: 0.6566 - accuracy: 0.7540  
Epoch 5/30  
86/86 [=====] - 94s 1s/step - loss: 0.4553 - accuracy: 0.8379  
Epoch 6/30  
86/86 [=====] - 93s 1s/step - loss: 0.2966 - accuracy: 0.9019  
Epoch 7/30  
86/86 [=====] - 93s 1s/step - loss: 0.1746 - accuracy: 0.9426  
Epoch 8/30  
86/86 [=====] - 93s 1s/step - loss: 0.1063 - accuracy: 0.9709  
Epoch 9/30  
86/86 [=====] - 93s 1s/step - loss: 0.0599 - accuracy: 0.9851  
Epoch 10/30  
86/86 [=====] - 93s 1s/step - loss: 0.0377 - accuracy: 0.9898  
Epoch 11/30  
86/86 [=====] - 92s 1s/step - loss: 0.0356 - accuracy: 0.9931
```

```
Epoch 12/30
86/86 [=====] - 92s 1s/step - loss: 0.0288 - accuracy: 0.99
24
Epoch 13/30
86/86 [=====] - 92s 1s/step - loss: 0.0227 - accuracy: 0.99
71
Epoch 14/30
86/86 [=====] - 92s 1s/step - loss: 0.0329 - accuracy: 0.99
09
Epoch 15/30
86/86 [=====] - 92s 1s/step - loss: 0.0355 - accuracy: 0.98
98
Epoch 16/30
86/86 [=====] - 92s 1s/step - loss: 0.0403 - accuracy: 0.98
87
Epoch 17/30
86/86 [=====] - 92s 1s/step - loss: 0.0495 - accuracy: 0.98
55
Epoch 18/30
86/86 [=====] - 92s 1s/step - loss: 0.0401 - accuracy: 0.98
69
Epoch 19/30
86/86 [=====] - 92s 1s/step - loss: 0.0527 - accuracy: 0.98
26
Epoch 20/30
86/86 [=====] - 92s 1s/step - loss: 0.0422 - accuracy: 0.98
80
Epoch 21/30
86/86 [=====] - 92s 1s/step - loss: 0.0094 - accuracy: 0.99
78
Epoch 22/30
86/86 [=====] - 92s 1s/step - loss: 0.0037 - accuracy: 0.99
96
Epoch 23/30
86/86 [=====] - 93s 1s/step - loss: 0.0029 - accuracy: 0.99
96
Epoch 24/30
86/86 [=====] - 92s 1s/step - loss: 0.0018 - accuracy: 0.99
96
Epoch 25/30
86/86 [=====] - 92s 1s/step - loss: 0.0036 - accuracy: 0.99
93
Epoch 26/30
86/86 [=====] - 92s 1s/step - loss: 0.0020 - accuracy: 0.99
96
Epoch 27/30
86/86 [=====] - 92s 1s/step - loss: 0.0043 - accuracy: 0.99
93
Epoch 28/30
86/86 [=====] - 93s 1s/step - loss: 0.0036 - accuracy: 0.99
93
Epoch 29/30
86/86 [=====] - 92s 1s/step - loss: 0.0025 - accuracy: 0.99
93
Epoch 30/30
86/86 [=====] - 92s 1s/step - loss: 0.0034 - accuracy: 0.99
93
Out[22]: <keras.callbacks.History at 0x1655680b130>
```

In [23]: `model.evaluate(X_test_scaled,y_test)`

```
29/29 [=====] - 44s 328ms/step - loss: 2.5548 - accuracy:
```

0.6558

Out[23]: [2.5548229217529297, 0.655773401260376]

Here we see that while train accuracy is very high (99%), the test accuracy is significantly low (66.99%) indicating overfitting. Let's make some predictions before we use data augmentation to address overfitting

In [24]:

```
predictions = model.predict(X_test_scaled)
predictions
```

Out[24]:

```
29/29 [=====] - 13s 314ms/step
array([[-8.034836, 25.752117, -1.1941547, -18.619875,
       -11.227778],
      [10.011834, -3.6381605, -11.387341, -5.632186,
       2.2027192],
      [-7.209098, 4.971569, 18.343872, -13.315592,
       -5.322657],
      ...,
      [3.8395944, -7.615803, -13.429502, 12.253704,
       9.793321],
      [4.1103244, -0.11814304, -9.373617, -3.000748,
       -0.36116636],
      [1.9852644, -13.491437, 0.55421954, 9.249021,
       7.549555]], dtype=float32)
```

In [25]:

```
score = tf.nn.softmax(predictions[0])
```

In [26]:

```
np.argmax(score)
```

Out[26]: 1

In [27]:

```
y_test[0]
```

Out[27]: 1

Improve Test Accuracy Using Data Augmentation

In []:

```
data_augmentation = keras.Sequential(
    [
        layers.experimental.preprocessing.RandomRotation(0.1),
        layers.experimental.preprocessing.RandomZoom(0.1),
    ]
)
```

Original Image

In []:

```
plt.axis('off')
plt.imshow(X[0])
```

Newly generated training sample using data augmentation

```
In [ ]: plt.axis('off')
plt.imshow(data_augmentation(X)[0].numpy().astype("uint8"))
```

Train the model using data augmentation and a drop out layer

```
In [ ]:
num_classes = 5

model = Sequential([
    data_augmentation,
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Dropout(0.2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])

model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

model.fit(X_train_scaled, y_train, epochs=30)
```

```
In [ ]: model.evaluate(X_test_scaled,y_test)
```

You can see that by using data augmentation and drop out layer the accuracy of test set predictions is increased to 73.74%

```
In [ ]:
```